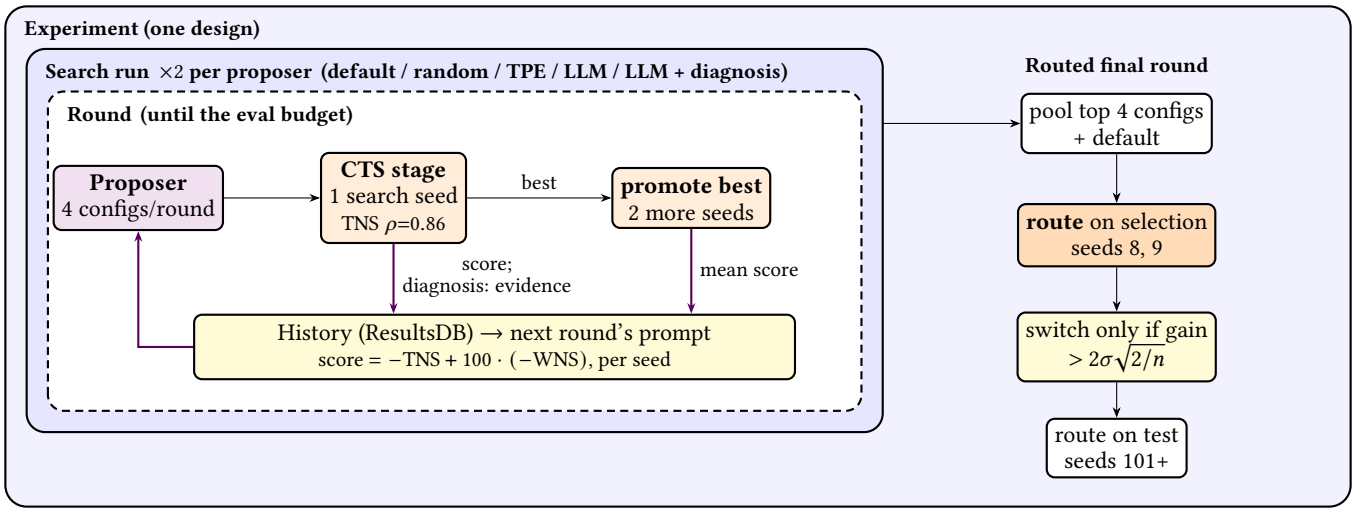

Experiment (one design)
Search run ×2 per proposer (default / random / TPE / LLM / LLM + diagnosis)
Round (until the eval budget)
Proposer
4 configs/round
CTS stage
1 search seed
TNS ρ=0.86
best
promote best
2 more seeds
score;
diagnosis: evidence
mean score
History (ResultsDB) → next round's prompt
score = −TNS + 100 · (−WNS), per seed
Routed final round
pool top 4 configs
+ default
route on selection
seeds 8, 9
switch only if gain
> 2σ√2/n
route on test
seeds 101+